SPECULA | Acquisition summary

11 November 2026 | Asia/Calcutta

Verified acquisitions with explicit remaining gates

The authoritative checklist has 136 tracked checkbox rows, including process and deliverable entries. Acquired resources include 9 research papers, 6 pinned repository clones, 5 pinned source ZIP archives, ATT&CK v19.2 STIX and official documentation. The file ledger records 2,254 acquired source/document files plus 4 preserved rejected responses (273.3 MiB total), including repository worktrees. This is not a claim that every checklist resource is complete.

First-ten pilot item	Status	Verified scope / remaining limit
1. NIST Hacking Case	DEFERRED / P0	Scenario, acquisition log and public answers acquired. Raw image 4.54 GiB deferred by owner; dependent experiments pending.
2. OpTC documentation	VERIFIED	README, eCAR, errata and red-team ground-truth PDF pinned to 5b108604f11f...; keys stored separately.
3. OpTC small events	UNAVAILABLE	Small public network request returned quota HTML. Endpoint host-group archive inspected is 1.86 GB and deferred.
4. DFIR-Metric release	PARTIAL	Three JSON datasets and generator notebooks acquired. Standalone evaluation runner absent at the pinned commit.
5. DFIR-Metric paper	VERIFIED	Actual arXiv 2505.19973v1 PDF downloaded, parsed and hashed.
6. ATT&CK STIX	VERIFIED	v19.2, 53,835,637 bytes, 26,086 objects; official Git blob hash matched. Existing XLSX exports preserved.
7. OCSF schema	VERIFIED	v1.9.0 source clone; version.json, exact commit and licence preserved. No installation.
8. NIST SP 800-86	VERIFIED	Official final PDF and publication source page present.
9. NIST SP 800-61r3	VERIFIED	Official April 2025 Revision 3 PDF and source page present.
10. Plaso / Timesketch	VERIFIED SCOPE	Both exact-commit sparse source/docs clones. Large test corpora omitted and recorded; neither installed.

Existing resources, integrity and limitations

Existing-resource audit

Inventoried 2,810 unique project files, including 322 source/fixture/config entries added to the primary inventory. Primary audit recorded 2,458 hashes and deferred 31 mutable datastore/index hashes. It identified 93 exact nonempty duplicate groups (654 member files); nothing was deleted. Named requested originals have a separate exact-name audit. Related Markdown/PDF files are not silently treated as the requested originals. Existing external or unseen files are not assumed absent.

Source and format verification

HTTP byte counts were reconciled where the source supplied Content-Length. Local SHA-256 values were recorded; repository commits and clean worktrees were checked with git connectivity verification. All acquired ZIP archives passed CRC checks; JSON parsed and PDFs passed strict structural/page/EOF checks. The original GenDFIR v4 PDF produces a non-zero-index xref warning in the PDF parser; its source bytes are preserved and the parser opens all pages. No publisher checksum is invented when one is not available.

Benchmark discrepancies

The actual DFIR-Metric release has 713 MCQ, 150 CTF and 510 forensic-search questions. Its README advertises 700/150/500. Exact source commit and actual counts are recorded; these are not independent label-quality validations. CFA-Bench supplies 20 cases; 59 raw PCAP/log files were copied with opaque filenames and byte-identical content. Original labels/results/name mappings remain in evaluator storage.

Bibliographic corrections

The corrected 2017 Drain paper (DOI 10.1109/ICWS.2017.13) and the separate 2018 DAG paper are acquired. CyberSleuth 2508.20643v2 is acquired with an author-linked sparse source clone. CFA-Bench DOI metadata and official benchmark are acquired; the publicly listed university author manuscript could be read through web tooling but validated direct downloads failed. No local CFA-Bench paper is claimed. MADIK DOI 10.1007/978-3-540-88875-8_10 remains a candidate needing relevance confirmation.

Access and size gates

Large evidence downloads remain deferred at the owner's request. NIST Hacking remains a P0 planned dataset, not completed or cancelled. NIST Data Leakage, CERT corpus releases, M57 images and large CyberSleuth/OpTC evidence are deferred. M57 instructor keys are encrypted/restricted; no access request or decryption was attempted. ISO source requests returned 403; UNB direct requests failed TLS validation; CERT small-file URLs returned 403. Zeek and Atomic Red Team source archives exceeded the 32 MiB automatic bound and were not saved. These failures are logged; no access checks were bypassed.

Evaluation boundary and reproducibility

Evaluator-only access

The evaluator directory has protected Windows ACL inheritance and grants DELL, SYSTEM and the acquisition sandbox account. This restricts other accounts but cannot isolate an evaluated agent sharing those identities. Enforce a distinct agent identity or evidence read-root policy before blinded evaluation. Give the system only the prepared CFA agent_evidence tree, after content-level leakage review; never expose inventories, original label-bearing filenames, tasks, results, original benchmark repos or mapping files. No existing parent answer keys were moved or repermissioned.

Templates, not research results

Empty templates exist for verified events, critical evidence, relationships, temporal order, original reviewer judgments, baseline/ablation measurements and resource logs. Reviewer instructions describe blinding and evidence-linked judgments. Experiment configuration, notebook and reproducibility manifest contain no invented labels, scores, model settings or runs. NIST-dependent experiments remain pending dataset acquisition and verification.

Deliverables

The master workbook contains Overview, Checklist, Acquired and Existing audit tabs. CSV manifests cover existing resources, exact-name checks, datasets, repositories, bibliography, acquired files, checksums, missing/pending and restricted resources. Acquisition receipts preserve official URLs, timestamps, HTTP headers, exact commits, sparse paths and source-archive versions. Preliminary report versions and the original checklist are retained.

Verification command

Run ..\venv\Scripts\python.exe 09_Acquisition_Reports\verify_campaign.py from the acquisition workspace as an authorized evaluator. It rechecks original and acquired hashes, original-to-evidence byte identity, repository status/commits, workbook counts/formula caches, manifest coverage, report structures and the unchanged parent source commit. Final PDF and workbook previews are inspected separately. Hashes establish observed bytes, not semantic truth or scientific suitability.

Next actions

Resolve the recorded access/source blockers without bypassing restrictions. Establish whether the MADIK candidate belongs to the original proposal. Select optional P2 resources according to a defined evaluation need and install tools only under separate approval. Obtain independent forensic labels and blinded reviewers; freeze models, budgets and protocol before experiments. Resume large evidence acquisition only under a new owner instruction.